1. From the Ethernet-attached machine with IP 10.11.6.1 (netmask 255.255.255.0), you use the command-line tool "ping" with argument 10.11.7.1. Supposing all caches are initially empty, list the work performed by each layer of the OSI stack, and the messages send up to the point where "ping" reports the first round-trip time.

   Let the sender be A, the router on the same LAN as A be R_A, the destination address be B and the router on the same network as B be R_B.

   Ping uses ICMP – which runs on top of IP. The method is "send a normal IP packet" with a payload corresponding to an ICMP packet. ICMP packets have the following form:

   

   Type indicates what type of ICMP message is being sent. This is 8 for ICMP echo (what Ping uses) and 0 for Echo reply (the response sent). Code is always zero. Data contains the first 8 bytes of the packet which caused the error – in the case of ping the premise is false so Data is empty. However, in general ICMP packets are sent after errors.

   The stages are:

   - A constructs the ICMP packet as described, with Type=8 and all other fields zero.
   - Next, A must send the IP packet to B.
   - A applies the netmask to the IP address of B and determines B is on a different LAN.
   - A *must* know its routers IP address (since it has an IP address). But it may not know its routers MAC address.
   - So it uses ARP to find its routers MAC address
   - A then builds the IP packet with the IP address corresponding to B and wraps this in a frame with destination MAC address for R_A. The frame is then sent to R_A.
   - R_A receives P. After buffering, R_A decrements the TTL – if it's zero, the packet is dropped. If nonzero the checksum is updated and the process continues. looks up $IP(B)$ in the forwarding table and finds no longest common prefix (since we assume no caches). So it initiates a routing protocol (ie Distance Vector Routing). The results of the routing protocol are stored in the Routing Table. This is then used to build a Forwarding Table – a mapping from routing prefix to port on which to forward the packet.

     Does the router decrement the TTL before or after buffering the packet? Just a small thing I can't see mentioned anywhere. I'll assume after.
   - R_A then looks up the IP address for B, sends P through the switching fabric to the corresponding output port. This routing/forwarding process is repeated at every router the packet encounters (although later routers shouldn't have to do routing before forwarding).
   - Eventually the packet reaches router R_B.
   - R_B performs ARP to find the address of B and forwards the packet onto B.
   - B sees it's a ping and responds with an ICMP echo.

- If the caches have been invalidated (or entries removed), the protocol is the same in reverse. However, ideally B would remember its routers MAC address, R_B would know the path to R_A and R_A would remember the MAC address for A – meaning the return journey would just be a number of lookups in forwarding tables.

ARP is implemented as follows:

- Each node keeps an ARP table which contains mappings from IP addresses to MAC addresses. This table is invalidated intermittently since IP addresses are not permanent.
- The node broadcasts a request on the LAN for the mapping from IP address to MAC address for node n.
- Each node receiving the message looks up the IP address in its ARP table. If there is an entry, it responds with the MAC address. If no entry is found, it broadcasts the request to all nodes (except the node the message came from).
- We assume the use of spanning tree or some other protocol to prevent forwarding loops.
- Every node on the path to the destination node now has an entry in its the forwarding table for the next node on the path. So the message can be sent.

Distance Vector Routing:

A distance vector D is a list of triples $(addressprefixes, subnetmask, distance, router)$ where $(a, s, d, r) \in D$ means that the subnet with address prefix a and subnet mask s can be reached in distance d by forwarding to router r.

- Each router continuously sends its distance vector (except the router to forward packets to) to its neighbours.
- On receiving a distance vector, the router sets its own distance vector to the elementwise minimum of the received distance vector (with the distance field incremented) and distance vector received (breaking ties deterministically by ID).
- The forwarding table is then constructed by adding the in the interface on which the packet should be sent to the distance vector.

2. What sequence of messages is required for a mobile device to lease an IP address using DHCP? What characteristic of a network requires the use of DHCP-Relay stations?

   There are four messages required for a mobile device A to lease an IP address using DHCP. Firstly, mobile device A broadcasts a message known as a DHCP discover with source IP address 0.0.0.0. The router hears this, if it has available IP addresses, it will respond directly to the client with a DHCP offer. Otherwise, if the router will first probe the network (broadcast "is anyone using this IP address") to find an IP address which is no longer in use.

   The mobile device A receives this DHCP offer and responds with a DHCP request (send directly to the router from the IP address it would like to use) – asking to use this IP address. This message is required since multiple routers may respond to the DHCP discover with different IP addresses. The router then responds with a DHCP Ack, acknowledging that the mobile device A is now registered as using that IP address.

   Routers have a range of IP addresses. If a node A wants to send a message to a node B, it will send a message to its router. However, if the node B is moving then it may not

   In DHCP relay, a router which isn't a DHCP server (DHCP relay agent) is able to forward DHCP messages to a DHCP server using its own IP address. The client sends

a DHCP message to the DHCP relay agent, which then forwards it to a DHCP server with its own IP address. The DHCP server then responds to the DHCP relay agent with the offer, which is forwarded to the client. The client responds to the DHCP relay agent with the request, which the DHCP agent forwards to the DHCP server. The server then responds with an acknowledgement which is forwarded to the client.

3. What does NAT do? Why do you have to configure a NAT gateway specifically if you want to play some network games on a home network?

   NAT maps public IP address / port combinations to ports on "private IP address". These are be used to address clients on the subnet. Each subnet can allocate IP addresses using conventional algorithms (ie DHCP). The NAT box then converts ports on private IP addresses to ports on a public IP address. This has two advantages. Firstly, a single IP address can be shared between thousands of devices. Secondly, NAT adds security: devices can be locally accessible but not publicly accessible (i.e my home printer is discoverable in my home but not on the wider internet).

   You have to configure a NAT gateway to make the private IP addresses on the home network globally visible – otherwise the external network would be unable to send messages to the user.

4. When do Ethernet switches need to run spanning tree protocol? How does it work?

   Forwarding loops occur when there is a loop $A \to B \to C \to \cdots \to A$ between switches $\{A, B, C, \dots\}$. When a switch A on a loop receives a frame destined for switch D which it cannot forward, it floods the network. If no switch in the loop has an entry for D in their forwarding table, they will all flood. Eventually, this frame will return to switch A. A still cannot forward to D so will flood the network again. This process repeats until a switch on the loop has a forwarding entry for switch D.

   Ethernet Switches need to run the Spanning Tree protocol when there are Ethernet links $A \to B \to C \to \cdots \to A$ between switches $\{A, B, C, \dots\}$. Note that this is only necessary when all devices on the loop are *switches* – since any device on the Transport Layer (router) would decrement the TTL in the IP header, leading to the packet eventually being discarded. Switches have no such guarantee. Spanning Tree is an algorithm to build a consistent spanning tree in a distributed manner. All links not on the spanning tree are disabled. This guarantees that every node is able to send a message to every other node, whilst also guaranteeing cycles are impossible.

   The Spanning Tree Protocol works as follows:

   - Each switch s broadcasts a message $(s, 0, s)$ "I am the root and I am distance 0 from myself".
   - On switch s_1 with state (r, d, s) receiving a message, (r', d', s_2) "switch s_2 believes switch r' is the root and is distance d' from it", switch s_1 will update its state to $\min\left((r, d, s) < (r', d' + 1, s_2)\right)$. If it changes its state, it will broadcast its new state to all its neighbours.

   The tree has been computed in a distributed manner (once its converged, nodes can tell their parents the link is open) and each node can now send messages to every other node with guarantees there are no

ICMP is a network layer protocol (since routers use it too!)

Ping begins by calling into the ICMP library.

You call a library and it fills in a byte array you pass to it.
Then pass it to another library which sends it (c.f. layering)

And with subnet masks in the forwarding table and see if the prefix addresses of any of the the interfaces on which you could send data match the destination.If so, use longest prefix match.
There is an entry 0.0.0.0/0 (which defaults to A router on the network)

ARP requires protocols: "from which protocol" and "to which protocol"
"Who has [IP 4 bytes, My IP address] [Ethernet 6 bytes, My MAC address] [IP 4 bytes, Dest IP address] [Ethernet 6 bytes, Dest MAC address]

You do NOT lookup ARP requests in your own ARP table -- let the authoritative source reply.
You just check whether the ARP request is for the interface on which you received that message.

Token passing gives guarantees -- so you need a specific bit of hardware to manage it.

Only fill in the ARP table entry if you get a message directly from that adapter

Ping calls a function.
This function does ARP (three retries 1s apart)
It THEN sends the ICMP packet
You then sleep until the correct ICMP packet returns.

There is no difference between a router and a host -- hosts can do forwarding (i.e mobile hotspot). They just usually have no good place to forward it to.

"Transport and Application layers are end-to-end"

DHCP is an Application Layer protocol -- the application layer sends to lower layers the IP address to send from and to.
It has to do a layering violation... it reads layer link layer adapters and says which adapter it wants to send out on.

Each layer has a field which indicates the protocol which is being run on the higher layer.

With DHCP, you send a message running port 67. This means "DHCP server". Every host on the LAN receives this message and wakes up all threads running protocol 67 (which is none on anything other than a DHCP server).

DHCP has a transaction ID -- just to ensure you're not duplicating any packets.

WE DO NOT USE ARP. WE DON'T USE ARP. THERE ARE 0 PACKETS WITH ETHERTYPE 0806.

DHCP works: you lease an IP address. You then sleep for half of this time and then wake up to renew the lease (with the same IP address).

Problem:
  DHCP sends a layer 2 broadcast. So it doesn't forward over WANs.
  Hence devices are physically incapable of getting an IP address if there is no DHCP server on the same ethernet as the sender...
DHCP relay resolves this. You have a router which forwards messages by unicast (broadcast doesn't exist on IP). These are manually configured.

In "normal NAT" (NAT without PAT), the NAT box translates traffic from one incoming internal IP address into one of a few outgoing IP addresses.
NAT without PAT allows you to have devices which are locally accessible.
Advantage of "NAT without PAT" compare to "NAT + PAT":
  You can setup connections with devices!
Problem:
  You're not increasing the number of devices which can simultaneously talk.

NAT + PAT:
  Address hosts with IP + port
  This increases the number of peers by 1000-65000x
Problem:
  Your port number is meaningless outside your network... this breaks FTP. There are plugins which fix this by scanning the binary and changing the port. This is a MASSIVE LAYERING VIOLATION. So you can't do FTP servers with TLS.

UDP multiplexes with source / destination ports and addresses.

IP, UDP, TCP don't say IP and port numbers are hidden. They are "exported". IP is designed such that every device is able to talk to every other device. NAT + PAT breaks that!

Ethernet was designed to be cheap and not particularly high efficiency. So they don't have TTL.
Using the switch with the lowest MAC address as the root was a mistake... because this is the oldest switch...

Spanning Tree:

Spanning Tree: "Bully Election"

Spanning Tree Protocol = STP. Later versions allow you to consider the link bandwidth.

Spanning tree takes n^2 messages in the worst case. You find the root and find the best path back to the root.
Always break ties with the lower number. 1st preference is path length and break ties with MAC address.

If you have multiple links between the same nodes, you tie break on the lower interface.
Bonding: "merging multiple links into one". You have to explicitly enable this... else you thrash the MAC address table.

Every 2s the root sends a 101 message "heartbeat".

If you don't hear from the root within 2s, the node restarts the protocol.

Ethernet error rate is tiny -- i.e. 1 bit in every 10^24 bits.

If a host hears a message from a node it's not aware of existing then it initiates spanning tree.

STP is always sent over every wire -- even wires which are turned off. Else if the real link were to fail you'd be unable to use the backup link to negotiate turning on the backup link.